

Algorithms Portfolio

Curated algorithmic solutions with correctness and complexity analysis

Ryan McMillan

MCompSci (cand.) – University of New England

C++ primary • Python secondary

February 28, 2026

License

This work is licensed under the MIT License. The canonical source and latest version are available at: <https://github.com/rmcmillan34/algorithms-portfolio>.

Glossary

Amortized analysis

An analysis technique that considers the average cost of operations over a sequence of operations, even if individual operations may be expensive.

Invariant

A property that remains true before and after each iteration of an algorithm or loop.

In-place algorithm

An algorithm that transforms its input using only a constant amount of additional memory.

Loop invariant

An invariant associated with a loop, used to prove correctness via initialization, maintenance, and termination.

Time complexity

An asymptotic measure of how an algorithm's running time grows as a function of input size.

Contents

Glossary	2
1 Introduction	4
1.1 How to Use This Book	4
1.2 Motivation	4
1.3 Acknowledgements	4
2 Arrays & Sequences	6
2.1 Core Ideas	6
2.2 Common Pitfalls	6
2.3 Primer - Arrays & Sequences	6
2.4 Easy	8
2.4.1 Plus One	8
2.4.2 Contains Duplicate	9
2.4.3 Single Number	11
2.4.4 Two Sum	12
3 Strings & Parsing	14
3.1 Core Ideas	14
3.2 Common Pitfalls	14
3.3 Easy	14
3.3.1 Valid Palindrome	14
4 Stacks and Queues	17
4.1 Core Ideas	17
4.2 Common Stack and Queue Variants	17
4.3 Common Pitfalls	17
4.4 Easy	17
4.4.1 Valid Parentheses	18
4.5 Medium	19
4.5.1 Add Two Numbers II	19
5 Linked Structures	23
5.1 Medium	23
5.1.1 Add Two Numbers	24

1 Introduction

This document is a curated portfolio of algorithmic solutions. Each selected problem includes:

- clear problem restatements,
- key invariants / observations,
- pseudocode,
- correctness sketches,
- complexity analysis,
- C++ (primary) and Python (secondary) implementations.

The focus is on clarity, correctness, and reasoning.

1.1 How to Use This Book

This book is not a tutorial on data structures and algorithms, nor does it attempt to replace formal coursework or foundational texts. Instead, it serves as a refresher and practical reference, organised around LeetCode-style problems that commonly appear in technical interviews and applied programming contexts.

The emphasis is on recognising patterns, recalling standard techniques, and reviewing concise, correct implementations. Explanations focus on the idea being exercised and the invariants maintained, rather than on full theoretical development.

Readers are assumed to already be familiar with fundamental data structures and algorithmic concepts. The book is intended to be used for revision, reinforcement, and interview preparation, particularly when revisiting topics after time away from formal study.

Problems include an optional **Author's Thinking** block, the authors plain language thinking and problem framing, which is excluded from upstream exemplar builds. Use this book as an example, to understand how to formally reason with, and structure algorithm and data structures problems.

1.2 Motivation

This portfolio was created as a structured, self-directed study of fundamental algorithms and data structures. While my formal coursework emphasizes applied computer science and systems-level engineering, it does not include a dedicated course focused on algorithmic problem-solving in the style commonly expected in technical interviews and quantitative software roles.

To address this, I undertook a systematic review of core algorithmic techniques, organizing problems by data structure and algorithmic paradigm, and documenting each solution with an explicit algorithm description, correctness argument, and complexity analysis. The goal is not only to arrive at correct implementations, but to practice communicating algorithmic reasoning clearly and rigorously.

This document serves both as a personal reference and as evidence of deliberate practice in algorithmic reasoning beyond formal coursework.

Many of the problems included here are sourced from commonly used interview problem sets, and were selected to emphasize invariant-based reasoning, edge case analysis, and asymptotic performance trade-offs.

Informal reasoning notes are included selectively to document thought process and learning progression; these sections are omitted in derived exemplar or presentation-focused versions of this material.

1.3 Acknowledgements

This work was informed by a number of open educational resources that emphasize rigorous algorithmic reasoning and clear exposition. In particular, the structure and approach to cor-

rectness arguments and asymptotic analysis were influenced by *MIT OpenCourseWare 6.006: Introduction to Algorithms*.

Additional reference was made to publicly available problem discussions, editorials, and textbooks where appropriate. All implementations, explanations, and written analyses in this document are my own, and any external influences are acknowledged at a conceptual level rather than reproduced directly.

2 Arrays & Sequences

Arrays are the foundational data structure in most algorithmic problems. They provide constant-time indexed access and strong cache locality, but require careful reasoning about indices, bounds, and in-place mutation.

This chapter focuses on recognizing array-based problem patterns and maintaining clear invariants over contiguous memory.

2.1 Core Ideas

- Index-as-state reasoning
- In-place mutation vs auxiliary storage
- Prefix and suffix accumulation
- Simulation of elementary operations (e.g. addition, carry propagation)

2.2 Common Pitfalls

- Off-by-one errors at boundaries
- Unintended data overwrites during in-place updates
- Assuming sorting or uniqueness without verifying constraints
- Over-allocating space when a single pass suffices

2.3 Primer - Arrays & Sequences

What It Is

sed.). An array is a *contiguous* block of memory storing elements of the same type, supporting $O(1)$ access by index. A dynamic array (e.g. C++ `std::vector`, Python `list`) grows by allocating a larger block and copying elements when capacity is exceeded, providing *amortized* $O(1)$ append. A *sequence* generalises the idea to any indexable or iterable ordered collection; here we emphasise contiguous arrays.

Recognition Patterns

- Index math across contiguous elements (prefix/suffix, window, ranges).
- Constraints emphasising in-place mutation, $O(1)$ extra space, single pass
- Repeated accumulation or difference of segments (prefix sums/scans).
- Local adjacency decisions: next/previous comparison, runs, peaks/valleys

Core Operations

```

1 access(A, i):
2     return A[i] // O(1)
3
4 traverse(A):
5     for i from 0 to n-1:
6         visit(A[i])
7
8 insert_at(A, i, x):
9     grow A by 1
10    for j from n-1 downto i:
11        A[j+1] = A[j]
12    A[i] = x
13
14 erase_at(A, i):
15    for j from i to n-2:
16        A[j] = A[j+1]
17    shrink A by 1
18
19 prefix_sums(A):
20    P[0] = 0
21    for i from 0 to n-1:
22        P[i+1] = P[i] + A[i]
23    return P
24
25 two_pointers_sorted(A, target):
26    i = 0; j = n - 1
27    while i < j:
28        s = A[i] + A[j]
29        if s == target: return (i, j)
30        if s < target: i = i + 1 else: j = j - 1
31    return not found
32
33 sliding_window(A):
34    l = 0
35    for r from 0 to n-1:
36        add(A[r])
37        while window_invalid():
38            remove(A[l])
39            l = l + 1
40    update_answer()

```

Listing 1: Arrays: Core Operations (pseudocode)

Invariants & Correctness

Array scans typically maintain index-based invariants: (i) after i steps, positions $[0..i)$ have been processed, (ii) two-pointer: indices satisfy $0 \leq i \leq j < n$ and a monotone move (left increases or right decreases) prevents rework, (iii) sliding window: the maintained sub-array $[l, r)$ satisfies a property (e.g. distinct elements, $\text{sum} \leq K$); expansion and contraction preserve the property.

Complexity

- Access: time $O(1)$; Traversal: $O(n)$
- Insert/delete at end (dynamic array): amortised $O(1)$, at index $O(n)$
- Prefix sums / scan: $O(n)$ time, $O(n)$ extra space (or $O(1)$ in place if possible)
- Two-pointers / sliding window: typically $O(n)$ if each pointer advances monotonically.

Common Pitfalls

- Off-by-one at boundaries; confusing inclusive/exclusive ranges.
 - Forgetting to update all window state on both expand and shrink
 - Overflow in sums/products; use wider types for accumulation.
 - Accidental $O(n^2)$ from repeated slicing/copying in inner loops.
 - Assuming sortedness/uniqueness without verifying constraints.
-)

2.4 Easy

The following problems introduce fundamental array manipulation techniques. The emphasis is on establishing and maintaining simple invariants while iterating through the array.

2.4.1 Plus One

Difficulty: Easy

Tags: Arrays, Simulation

ID: LC-066

Problem. Given a non-empty array of digits representing a non-negative integer, increment the integer by one and return the resulting array of digits.

The most significant digit is stored at the beginning of the array, and each digit is in the range $[0, 9]$.

Key idea. Incrementing the number requires simulating elementary addition from right to left. Carry propagation only affects trailing digits equal to nine; once a digit strictly less than nine is incremented, the process terminates.

Author's Thinking.

My initial instinct was to convert the digit array into an integer, add one, and convert it back. This approach fails due to potential overflow and violates the problem constraints.

What simplified the reasoning was treating the operation as grade-school addition and maintaining the invariant that all digits to the right of the current index are already correct after carry handling.

Algorithm.

```
1 for i from n - 1 downto 0:
2     if digits[i] < 9:
3         digits[i] = digits[i] + 1
4         return digits
5     digits[i] = 0
6
7 prepend 1 to digits
8 return digits
```

Listing 2: Pseudocode

Correctness sketch. We iterate from the least significant digit toward the most significant digit. At each iteration, the invariant maintained is that all digits to the right of the current index represent the correct result after increment and carry propagation.

If a digit less than nine is encountered, incrementing it produces the correct result and no further carry is required. If all digits are nine, each is set to zero and a leading one is prepended, yielding the correct incremented number. Thus, the algorithm always returns the correct result.

Complexity. The algorithm runs in $O(n)$ time, where n is the number of digits. It uses $O(1)$ additional space beyond the output array.

C++ Implementation

```
1  /*
2  * LC-066 - Plus One
3  *
4  */
5
6  #include <iostream>
7
8  class Solution {
9  public:
10     vector<int> plusOne(vector<int>& digits) {
11         for(int i = digits.size(); i-- > 0 ; ) {
12             if (digits[i] != 9) {
13                 digits[i]++;
14                 return digits;
15             }
16             else {
17                 digits[i] = 0;
18             }
19         }
20
21         vector<int> newDigits(digits.size() + 1);
22         newDigits[0] = 1;
23         for (int j = 1; j < newDigits.size(); j++) {
24             newDigits[j] = digits[j-1];
25         }
26
27         return newDigits;
28     }
29 };
```

Listing 3: C++ Implementation

Python Implementation

```
1  # LC-066 Plus One
2
3  from typing import List
4
5  def plusOne(digits: List[int]) -> List[int]:
6      for i in range(len(digits) - 1, -1, -1):
7          if digits[i] < 9:
8              digits[i] += 1
9              return digits
10         digits[i] = 0
11
12     return [1] + digits
```

Listing 4: Python Implementation

2.4.2 Contains Duplicate

Difficulty: Easy

Tags: Arrays, Hash Table, Sorting

ID: LC-217

Problem. Given an integer array *nums*, return true if any value appears at least twice in the array, and return false if every element is distinct.

Key idea. Iterate over the array, track and compare which integers have been seen before.

Author's Thinking.

The brute force method that I think of initially is create a new array, for every number I see in *nums*, check if it is in the new array, if not append it to the new array, then move to the next *nums[i]*. This results in having to iterate through two arrays for every iteration.

We can make the above better by implementing a *set* allowing for $O(1)$ lookup of seen integers.

Algorithm.

```
1  let seen be a set
2
3  for num in nums:
4      if num in seen:
5          return true
6      else:
7          insert num into hm
8
9  return false
```

Listing 5: Pseudocode

Correctness sketch. As we are iterating from the first element, to the final element of the array; for every iteration, we check if the integer is in the hash map and return true early before insertion. we know that all elements to the left of our index processed and can be certain have not appeared in the array as a duplicate.

Complexity. The algorithm runs in $O(n)$ time, where n is the number of digits in the input array. It has an auxiliary space complexity of $O(n)$

C++ Implementation

```
1  class Solution {
2  public:
3      bool containsDuplicate(vector<int>& nums) {
4          std::unordered_set<int> seen;
5          for (int num : nums) {
6              // Insert returns a std::pair<iterator, bool> where .first is iterator
7              if (!seen.insert(num).second) return true;
8          }
9          return false;
10     }
11 };
```

Listing 6: C++ Implementation

Python Implementation

```
1  from typing import List
2
3  class Solution:
4      def containsDuplicate(self, nums: List[int]) -> bool:
5          seen = set()
```

```

6     for x in nums:
7         if x in seen:
8             return True
9         seen.add(x)
10    return False

```

Listing 7: Python Implementation

2.4.3 Single Number

Difficulty: Easy

Tags: Arrays, Bitwise Operations

ID: LC-136

Problem. TODO

Key idea. TODO

Author's Thinking.

TODO

Algorithm.

```

1  TODO

```

Listing 8: Pseudocode

Correctness sketch. TODO

Complexity. TODO

C++ Implementation

```

1  class Solution {
2      public:
3          int singleNumber(vector<int>& nums) {
4              int ans = 0;
5              // Perform bitwise XOR (ans XOR nums[i])
6              for(int x : nums) ans ^= x;
7              return ans;
8          }
9  }

```

Listing 9: C++ Implementation

Python Implementation

```

1  class Solution:
2      def singleNumber(self, nums: List[int]) -> int:
3          ans = 0
4          for x in nums:
5              ans = x ^ ans
6          return ans

```

Listing 10: Python Implementation

2.4.4 Two Sum

Problem. Given an array of integers *nums* and an integer *target*, return the *indices* of the two numbers such that they add up to *target*.

You may assume that each input would have **exactly one solution**, and you may not use the same element twice.

You can return the answer in any order.

Key idea. We must find complementary numbers that sum to the given target integer. By using a hash table, we can store already-seen numbers and check if their complement (the difference between the target and the current number) has been encountered previously. This allows constant-time lookup for finding the solution.

Author's Thinking.

My initial thoughts for this problem included a brute-force solution iterating over the entire input array. For each integer (*i*) in the array we can compute the complement by $target - i = complement$ and then continue iterating over the array until we find a match noting there is exactly one solution. This would result in a time complexity of $O(n^2)$. A better solution would be to iterate over the array once and use constant time lookup using a hash table.

Algorithm.

```

1  hash_map<complement, index> hm;
2
3  for num in input:
4      complement = target - num
5      if not complement in hm
6          hm[num] = index
7      return [index, hm[complement]]
8  else:
9      return [index, hm[complement]]

```

Listing 11: Pseudocode

Correctness sketch. We know that the algorithm is correct as we iterate over the input array and check if the complement exists in the hash map for every iteration. The invariant that we maintain is that we know that every integer that has already been seen's complement does not yet exist in the hash map. Given the problem statement condition that we can assume *exactly one solution*, this allows us to return the answer immediately.

Complexity. let n be the length of the input array. For the hash map implementation we iterate over the input array a single time, and each access of the hash map is a constant time operation. therefore, the time complexity of this algorithm is $O(n)$. The auxiliary space complexity in the worst case i.e. where we find the complementary integer last means that we store all n integers in the hash map, making the auxiliary space complexity is also $O(n)$.

C++ Implementation

```

1  /*
2  * LC-001 - Two Sum
3  *
4  */
5
6  class Solution {

```

```
7 public:
8     vector<int> twoSum(vector<int>& nums, int target) {
9         std::unordered_map<int, int> seen;
10
11         for (int i=0; i<nums.size(); i++) {
12             int complement = target - nums[i];
13
14             if (!seen.contains(complement)) {
15                 seen.insert({nums[i], i});
16             }
17             else {
18                 return std::vector {i, seen[complement]};
19             }
20         }
21         return {};
22     }
23 };
```

Listing 12: C++ Solution: Two Sum

Python Implementation

```
1 # LC-001 Two Sum
2
3 from typing import Sequence
4 from typing import List
5
6 class Solution:
7     def twoSum(self, nums: List[int], target: int) -> List[int]:
8         seen = {}
9
10        for idx, num in enumerate(nums):
11            complement = target - num
12            if not complement in seen:
13                seen[num] = idx
14            else:
15                return [idx, seen[complement]]
```

Listing 13: Python Solution: Two Sum

3 Strings & Parsing

Strings are a fundamental data structure used to represent text and sequences of characters. They often require specialized algorithms for efficient manipulation, searching, and pattern matching. This chapter focuses on recognizing string-based problem patterns and maintaining clear invariants over character sequences.

3.1 Core Ideas

- Two-pointer techniques for substring problems
- Sliding window approaches for dynamic substring analysis
- Character classification and normalisation
- Maintaining loop invariants over indices
- Incremental Parsing
- In-place vs Auxiliary Storage

3.2 Common Pitfalls

- Off-by-one errors at boundaries
- Incorrect handling of empty strings or single-character strings
- Failing to advance pointers correctly when skipping characters
- Unnecessary string copying leading to increased time/space complexity
- Assumptions about character sets (e.g., ASCII vs Unicode)

3.3 Easy

The following problems introduce fundamental string parsing and manipulation techniques. The emphasis is on establishing and maintaining simple invariants while iterating through the array.

3.3.1 Valid Palindrome

Difficulty: Easy

Tags: Strings, Two Pointers

ID: LC-125

Problem. Given a string s , determine whether it is a palindrome when considering only alphanumeric characters and ignoring case. Non-alphanumeric characters (punctuation/whitespace) are ignored. After filtering, the empty string is considered a palindrome.

Key idea. We need to compare the string from both ends; a two-pointer approach lets us scan inward while skipping non-alphanumeric characters, achieving $O(n)$ time and $O(1)$ extra space.

Author's Thinking.

<Informal reasoning, failed approaches, intuition-building, or mental models. This section is intentionally exploratory and may be omitted in promoted or exemplar versions of the book.>

Algorithm.

```

1  initialise left = 0, right = length(s) - 1
2  while left < right:
3      if s[left] is not alphanumeric:
4          left = left + 1
5          continue
6      if s[right] is not alphanumeric:
7          right = right - 1
8          continue
9      if lowercase(s[left]) != lowercase(s[right]):
10         return false
11     left = left + 1
12     right = right - 1
13 return true

```

Listing 14: Pseudocode

Correctness sketch. At the start of each iteration, all alphanumeric characters strictly outside the interval $[left, right]$ have already been matched (case-insensitively) with their corresponding partner on the other side. Skipping a non-alphanumeric character is safe because such characters are ignored by the problem definition. If $\text{lower}(s[left]) \neq \text{lower}(s[right])$, then an unmatched pair exists, so the string cannot be a valid palindrome and we return **False**. When the loop terminates ($left \geq right$), every relevant pair has been verified; a middle character (odd length after filtering) does not affect palindromicity.

Complexity. Each pointer moves inward at most n times, so time complexity is $O(n)$. The algorithm uses $O(1)$ extra space (in-place scanning; no filtered string constructed).

Python Implementation

```

1  # LC 125. Valid Palindrome
2
3  class Solution:
4      def isPalindrome(self, s: str) -> bool:
5          """
6          Determines if a given string is a valid palindrome, considering only
7          alphanumeric characters and ignoring case sensitivity.
8
9          Args:
10             s (str): The input string.
11
12          Returns:
13             bool: True if the string is a palindrome, False otherwise.
14          """
15         # Initialise two pointers
16         left = 0
17         right = len(s)-1
18
19         # Iterate until the left and right pointers meet
20         while (left < right):
21             if not s[left].isalnum():
22                 left += 1
23                 continue
24             if not s[right].isalnum():
25                 right -= 1
26                 continue
27             if s[left].lower() != s[right].lower():
28                 return False

```

```
29
30     left += 1
31     right -= 1
32
33     return True
```

Listing 15: Python Solution: LC-125 Valid Palindrome

4 Stacks and Queues

Stacks and queues are fundamental abstract data types that model **sequential access patterns**. They are often implemented using arrays or linked lists, but their defining characteristic is the order in which elements are accessed and modified.

Stacks follow a **Last-In-First-Out (LIFO)** order, where the most recently added element is the first to be removed. Queues follow a **First-In-First-Out (FIFO)** order, where the earliest added element is the first to be removed. These structures are essential for a wide range of algorithms, including depth-first search (DFS), breadth-first search (BFS), and various parsing techniques.

4.1 Core Ideas

- **Access-order invariant:** only interact with one end (stack top / queue front), which simplifies reasoning and correctness.
- **Monotonic structures:** maintain increasing/decreasing order to answer "next greater/smaller" or "sliding window max/min" in linear time.
- **State simulation:** use stacks/queues to simulate recursion, parsing, or reverse traversal (e.g., LC-445 uses LIFO to process from least-significant digit).

4.2 Common Stack and Queue Variants

- **Monotonic Stack** — Maintains increasing/decreasing values (often indices) to solve Next Greater/Smaller patterns.
- **Monotonic Deque** — Maintains candidates for window max/min; front is the answer, back removes dominated elements.
- **Circular Queue** — Fixed-capacity queue implemented with two pointers (common in design questions).
- **Two-Stack Queue / Two-Queue Stack** — Classic ADT transformations; highlights amortized analysis.
- **Deque** — Double-ended queue enabling efficient push/pop at both ends (basis for sliding window problems).

4.3 Common Pitfalls

- **Empty checks and order of operations:** calling `top()`/`front()`/`back()` before verifying non-empty.
- **Value vs index confusion:** many patterns must store *indices* (for window bounds, distance, duplicates) not values.
- **Monotonic maintenance mistakes:** using the wrong inequality ($>$, $>=$) causing duplicate-handling bugs.
- **Queue window eviction errors:** forgetting to evict out-of-window indices from the front (off-by-one boundaries).
- **Assuming `pop()` returns a value:** e.g., C++ `vector::pop_back()` / `stack::pop()` return void; read via `back()`/`top()` first.

4.4 Easy

The following problems build foundational stack/queue instincts. The emphasis is on the **access-order invariant** (LIFO/FIFO), clean **empty-check discipline** before `top()`/`front()`, and using a stack to model **unmatched state** while scanning left-to-right (common in parsing).

4.4.1 Valid Parentheses

Difficulty: Easy

Tags: Stacks

ID: LC-020

Problem. Given a string containing only the characters `()[]{}` , determine whether it is valid. A string is valid if every opening bracket is closed by the same type of bracket, and brackets are closed in the correct order.

Key idea. Use a stack to represent the sequence of currently-unmatched opening brackets. When an opening bracket is seen, push it. When a closing bracket is seen, it must match the most recent unmatched opening bracket (the stack top). Any mismatch, or attempting to close when the stack is empty, makes the string invalid. The string is valid iff the stack is empty at the end.

Author's Thinking.

My first attempt jumped into coding the stack loop before writing the invariant. I was conceptually correct (LIFO matching), but I missed two edge cases that matter in C++: (1) calling `top()` on an empty stack (crash/UB), and (2) ensuring we `pop()` after a successful match (easy to accidentally skip with `continue`). Writing the invariant first made the implementation obvious: “stack = exactly the unmatched openers so far”.

Algorithm.

```

1 stack = empty
2 for c in s:
3     if c is '(' or '[' or '{':
4         push c onto stack
5     else:
6         if stack is empty: return false
7         t = top of stack
8         if (c,t) is not one of (')', '}', ']', '}', '}', '{'):
9             return false
10    pop stack
11 return (stack is empty)

```

Listing 16: Pseudocode

Correctness sketch. Maintain the invariant: after processing any prefix of the string, the stack contains exactly the opening brackets from that prefix that have not yet been matched, in the order they were encountered. When an opening bracket is read, pushing it preserves the invariant by adding a new unmatched opener. When a closing bracket is read, validity requires that there exists an unmatched opener and that the most recent unmatched opener matches the closing type; this is exactly the stack top. If the stack is empty or the types mismatch, the string cannot be valid, and the algorithm returns false. Otherwise, popping removes the matched opener and preserves the invariant. After the full scan, the string is valid if no unmatched openers remain, i.e., the stack is empty. Therefore, the algorithm returns true exactly for valid strings.

Complexity. Let n be the length of the string. Each character is pushed and popped at most once, so the runtime is $O(n)$. The stack stores at most n opening brackets, so the extra space is $O(n)$.

C++ Implementation

```

1 #include <bits/stdc++.h>
2

```

```
3 class Solution {
4 public:
5     bool isValid(string s) {
6         std::stack<char> brackets;
7         std::string::iterator it;
8
9         for (it = s.begin(); it != s.end(); ++it) {
10             char c = *it;
11             if (*it == '(' || *it == '[' || *it == '{') {
12                 brackets.push(*it);
13             }
14             else {
15                 if ( brackets.empty() ) return false;
16
17                 char current = brackets.top();
18                 brackets.pop();
19
20                 if ( current == '{' && c != '}' ) return false;
21                 if ( current == '[' && c != ']' ) return false;
22                 if ( current == '(' && c != ')' ) return false;
23             }
24         }
25         return brackets.empty();
26     }
27 };
```

Listing 17: C++ Solution: Valid Parentheses

Python Implementation

```
1 class Solution:
2     def isValid(self, s: str) -> bool:
3         st = []
4         match = {')': '(', ')': '[', ']': '{'}
5
6         for c in s:
7             if c in "([{":
8                 st.append(c)
9             else:
10                 if not st or st[-1] != match.get(c):
11                     return False
12                 st.pop()
13
14         return not st
```

Listing 18: Python Solution: Valid Parentheses

4.5 Medium

The following problems build on basic stack and queue traversal and introduce **stateful pointer algorithms**, where correctness depends on maintaining multiple interacting invariants simultaneously.

4.5.1 Add Two Numbers II

Difficulty: Medium

Tags: Linked Lists, Math, Stack

ID: LC-445

Problem. Given two non-empty linked lists representing two non-negative integers, where each node stores a single digit and the digits are arranged in normal order (most significant digit first), compute the sum of the two integers. Return the result as a new linked list in the same format.

Key idea. TODO

Author's Thinking.

LC-445 feels like LC-002, except the digits are stored in forward order (most-significant first). That breaks the usual “add while traversing” approach, because the carry depends on digits at the end of the list that we haven’t reached yet. Reversing both lists would reduce it to LC-002, but that either mutates the inputs or adds extra reversal steps.

Instead, I used two stacks, leveraging their LIFO (Last-In, First-Out) behavior to simulate processing the lists from right to left. I traverse each list once, pushing digits onto a stack; then popping gives the digits back in reverse order (least-significant first), which makes addition with carry straightforward.

As I compute each digit, I prepend it to the result list so the final output remains in forward order without needing to reverse anything.

Algorithm.

```

1  s1, s2 = empty stacks
2
3  for node in l1:
4      push node.val onto s1
5  for node in l2:
6      push node.val onto s2
7
8  carry = 0
9  head = null
10
11 while s1 not empty OR s2 not empty OR carry != 0:
12     x = pop(s1) if s1 not empty else 0
13     y = pop(s2) if s2 not empty else 0
14
15     total = x + y + carry
16     carry = total // 10
17     digit = total % 10
18
19     // prepend digit to result
20     head = new ListNode(digit, head)
21
22 return head

```

Listing 19: Pseudocode

Correctness sketch. TODO

Complexity. TODO

C++ Implementation

```

1  /** LC-445
2   * Definition for singly-linked list.
3   * struct ListNode {
4   *     int val;
5   *     ListNode *next;

```

```

6  * ListNode() : val(0), next(nullptr) {}
7  * ListNode(int x) : val(x), next(nullptr) {}
8  * ListNode(int x, ListNode *next) : val(x), next(next) {}
9  * };
10 */
11 class Solution {
12 public:
13     ListNode* addTwoNumbers(ListNode* l1, ListNode* l2) {
14         // Build Stacks s1,s2 from l1,l2.
15         std::vector<int> s1, s2;
16
17         while (l1) { s1.push_back(l1->val); l1 = l1->next; }
18         while (l2) { s2.push_back(l2->val); l2 = l2->next; }
19
20         // Perform arithmetic on stack values
21         int carry = 0;
22         ListNode* head = nullptr;
23
24         while ( !s1.empty() || !s2.empty() || carry ){
25             int x = 0, y = 0;
26             if ( !s1.empty() ) { x = s1.back(); s1.pop_back(); }
27             if ( !s2.empty() ) { y = s2.back(); s2.pop_back(); }
28
29             int total = x + y + carry;
30             int digit = total % 10;
31             carry = total / 10;
32
33             // Prepend to build the result in forward order (since we pop from the
34             // back).
35             ListNode* node = new ListNode(digit);
36             node->next = head;
37             head = node;
38         }
39         return head;
40     }
41 };

```

Listing 20: C++ Solution: LC-445 Add Two Numbers II

Python Implementation

```

1  # LC-445
2  # Definition for singly-linked list.
3  # class ListNode:
4  #     def __init__(self, val=0, next=None):
5  #         self.val = val
6  #         self.next = next
7
8  class Solution:
9      def addTwoNumbers(self, l1: ListNode, l2: ListNode) -> ListNode:
10         s1, s2 = [], []
11
12         while l1:
13             s1.append(l1.val)
14             l1 = l1.next
15         while l2:
16             s2.append(l2.val)

```

```
17     l2 = l2.next
18
19     carry = 0
20     head = None
21
22     while s1 or s2 or carry:
23         x = s1.pop() if s1 else 0
24         y = s2.pop() if s2 else 0
25
26         total = x + y + carry
27         carry = total // 10
28
29         # Prepend to keep forward order (we're adding from least to most
30         significant).
31         head = ListNode(total % 10, head)
32
33     return head
```

Listing 21: Python Solution: LC-445 Add Two Numbers II

5 Linked Structures

Linked structures are **pointer-based** data structures where logical order is defined by explicit **references** rather than **contiguous memory** layout. They sacrifice constant-time indexed access in exchange for flexible insertion, deletion, and natural modeling of sequential processes.

This chapter focuses on recognising pointer-driven problem patterns and maintaining correct **traversal invariants** over **node-based** structures.

Core Ideas

- **Pointer-as-state** reasoning
- Sequential traversal instead of indexed access
- Explicit handling of **null termination**
- Incremental construction using moving **ursors** (e.g. **dummy head** nodes)

Common Linked Structure Variants

Linked structures appear in several common forms, each with distinct properties and use cases:

- **Singly Linked List** — each node references the next node only
- **Doubly Linked List** — each node references both next and previous nodes
- **Circular Linked List** — the terminal node references an earlier node
- **Sentinel-Based List** — uses a non-data **sentinel node** to simplify edge cases

Common Pitfalls

- Losing access to the ****head node**** during traversal or mutation
- Dereferencing ****null pointers****
- Incorrect ****loop termination conditions****
- Conflating ****in-place modification**** with ****new list construction****

Choosing Between Linked Structures & Arrays

Selecting the appropriate data structure is often as important as implementing the algorithm itself. Arrays and Linked Structures optimise for different access and update patterns.

Prefer Linked Structures When:

- Order matters, but **indexing does not**,
- Frequent insertions happen towards the end,
- Data arrives or is processed **sequentially**,
- The problem naturally models as pointer traversal,
- The input is provided as a **node reference** rather than a collection.

Key Trade Offs

- Linked Structures provide $O(1)$ insertion, $O(n)$ access
- Arrays provide $O(1)$ access, $O(n)$ insertion

5.1 Medium

The following problems build on basic linked-structure traversal and introduce **stateful pointer algorithms**, where correctness depends on maintaining multiple interacting invariants simultaneously. Rather than simple scans, these problems require coordinating traversal with auxiliary

state (such as a **carry**, window offset, or delayed reference) while constructing or modifying linked structures.

The emphasis is on:

- Maintaining multiple **traversal invariants**
- Coordinating pointer movement with additional state
- Incremental construction using **dummy heads** and **cursors**
- Correct handling of **null termination** across uneven structures

5.1.1 Add Two Numbers

Difficulty: Medium

Tags: Linked Lists, Math, Carry Propagation

ID: LC-002

Problem. Given two non-empty linked lists representing two non-negative integers, where each node stores a single digit and the digits are arranged in reverse order (least significant digit first), compute the sum of the two integers. Return the result as a new linked list in the same reverse-order format.

Key idea. Because the digits are stored in reverse order, the lists can be traversed from least significant to most significant digit, making the problem a direct simulation of elementary addition. At each step, the current digits and any incoming **carry** determine the next output digit and updated carry. A **dummy head** node anchors the result list, while a moving **cursor** appends newly computed digits without losing access to the list's beginning. Traversal continues until both input lists reach **null termination** and no carry remains, maintaining the **invariant** that all previously processed digits in the result list are correct.

Author's Thinking.

A first brute-force idea is to traverse each list, reconstruct the two integers, add them, and then convert the sum back into a linked list. However, this approach is awkward in practice: the numbers may exceed native integer ranges, and it performs extra passes and conversions that are unnecessary.

A better approach is to simulate column-wise addition directly on the linked lists. Since the digits are stored least-significant-first, we can traverse both lists once, propagate a carry, and build the result list incrementally.

Algorithm.

```
1  dummy = new Node(0)
2  cur = dummy
3  carry = 0
4
5  while l1 != null OR l2 != null OR carry != 0:
6      x = (l1 != null) ? l1.val : 0
7      y = (l2 != null) ? l2.val : 0
8
9      total = x + y + carry
10     digit = total mod 10
11     carry = total div 10
12
13     cur.next = new Node(digit)
14     cur = cur.next
15
16     if l1 != null: l1 = l1.next
17     if l2 != null: l2 = l2.next
18
19  return dummy.next
```

Listing 22: Pseudocode

Correctness sketch. At the start of each iteration of the loop, the result list contains the correct sum of all digit positions processed so far, and the **carry** variable holds the correct carry-out into the next higher digit position. During the iteration, the algorithm adds the current digits (or zero if a list has reached **null termination**) together with the carry, computes the correct output digit, and updates the carry accordingly. Because each iteration advances at least one input pointer or consumes a remaining carry, the loop eventually terminates. When the loop ends, all digit positions and any final carry have been processed, so the resulting linked list represents the correct sum.

Complexity. Let n and m be the lengths of the two input linked lists. Each iteration processes at most one node from each list, and the loop runs for at most $\max(n, m) + 1$ iterations, yielding a time complexity of $O(\max(n, m))$.

The algorithm uses $O(1)$ auxiliary space, as it maintains only a constant number of pointers and a carry variable. The output linked list requires $O(\max(n, m))$ space to store the result.

C++ Implementation

```

1  /**
2   * Definition for singly-linked list.
3   * struct ListNode {
4   *   int val;
5   *   ListNode *next;
6   *   ListNode() : val(0), next(nullptr) {}
7   *   ListNode(int x) : val(x), next(nullptr) {}
8   *   ListNode(int x, ListNode *next) : val(x), next(next) {}
9   * };
10  */
11
12  class Solution {
13  public:
14      ListNode* addTwoNumbers(ListNode* l1, ListNode* l2) {
15          // Create two ListNode pointers to track the results beginning and current
16          // address.
17          ListNode* dummy = new ListNode(0);
18          ListNode* current = dummy;
19
20          // Initialise our carry variable to 0
21          int carry = 0;
22
23          // Iterate while we have a non null pointer in L1 or L2 or a carry digit
24          while (l1 || l2 || carry) {
25              // If Ls contain value use that, otherwise 0
26              int x = l1 ? l1->val : 0;
27              int y = l2 ? l2->val : 0;
28
29              // Perform calculation for result digit and carry
30              int total = x + y + carry;
31              carry = total / 10;
32              int digit = total % 10;
33
34              // Create memory for a new result node, and point to that node
35              current->next = new ListNode(digit);
36              current = current->next;

```

```
36         // Traverse the linked lists if there is a valid node
37         if (l1) l1 = l1->next;
38         if (l2) l2 = l2->next;
39
40     }
41
42     return dummy->next;
43
44 }
45 };
46
```

Listing 23: C++ Solution: LC-002 Add Two Numbers

Python Implementation

```
1  # Definition for singly-linked list.
2  # class ListNode:
3  #     def __init__(self, val=0, next=None):
4  #         self.val = val
5  #         self.next = next
6
7  class Solution:
8      def addTwoNumbers(self, l1: Optional[ListNode], l2: Optional[ListNode]) ->
9          Optional[ListNode]:
10         # Initialise output LinkedList, and carryover integer
11         current = ListNode(0)
12         dummy = current
13         carry = 0
14
15         # Traverse both lists while digits or carry remain
16         while(l1 or l2 or carry):
17             x = l1.val if l1 else 0
18             y = l2.val if l2 else 0
19
20             total = x + y + carry
21             digit = total % 10
22             carry = total // 10
23
24             # Append current result digit and advance
25             current.next = ListNode(digit)
26             current = current.next
27
28             l1 = l1.next if l1 else None
29             l2 = l2.next if l2 else None
30
31         return dummy.next
```

Listing 24: Python Solution: LC-002 Add Two Numbers